

## DENORMALIZATION:

One of the advantages of using MongoDB over a relational database like MySQL is that unstructured data can be stored in MongoDB. This along with the scalability facilitated by the embedded document structure of data and scale-out architecture results in low effort to set up, unlike MySQL which requires thorough schema design with primary keys, foreign keys, constraints and so on. This advantage of MongoDB allows us to design denormalized databases. Denormalization is based on the simple principle of *“Data that is accessed together should be stored together”*. Denormalized databases can improve read performance and query performance in a variety of cases, such as:

- A recurring query requires a few fields from a large document in another collection. We can choose to maintain a copy of those fields in an embedded document in the collection that the recurring query targets to avoid merging two distinct collections or performing frequent \$lookup operations
- An average value of some field in a collection is frequently requested. We can choose to create a derived field in a separate collection that is updated as part of your writes and maintains a running average for that field

While embedding documents or arrays without data duplication is preferred for grouping related data, denormalization can improve read performance when separate collections must be maintained. To summarize, denormalization makes sense when we have a high read-to-write ratio [\[Link to official guide\]](#).

In the case of our database, unlike our previous proposed architecture in Lab 1 wherein we had designed an elaborate ER diagram for a relational database with detailed relationships between the entities (i.e. tables), in the MongoDB database case, we can use the principle of denormalization to consolidate multiple tables into 3 collections:

- Companies collection –
  - Combination of the companies, company\_industries and company\_specialties datasets, all joined together by company\_id
  - Industries and specialties are embedded into arrays to reduce redundancy
  - Snapshots of Python code used to create the underlying dataframe for companies collection below:

# Grouping industries df by company\_id and converting industry into a list

```
company_industries_grouped_df = company_industries_df.groupby('company_id').agg({'industry':  
lambda x: list(x)}).reset_index()
```

# Grouping specialties df by company\_id and converting specialties into a list

```
company_specialties_grouped_df = company_specialties_df.groupby('company_id').agg({'specialty':  
lambda x: list(x)}).reset_index()
```

# Merging all datasets with companies\_df being the left table

```
merged_companies_df = companies_df.merge(company_industries_grouped_df, on='company_id',
how='left').merge(company_specialities_grouped_df,on='company_id',how='left')
```

```
# Renaming the column name to company_name for better understanding
```

```
merged_companies_df.rename(columns = {'name':'company_name'}, inplace = True)
```

- Job postings –
  - This collection combines the job\_postings, job skills, job\_benefits and job\_industries datasets all joined together by job\_id
  - Skills, benefits and industries are embedded into arrays to reduce redundancy
  - Snapshots of Python code used to create the underlying dataframe for job\_postings collection below:

```
# Grouping benefits df by job_id and converting benefits (type column) into a list
```

```
benefits_grouped_df = benefits_df.groupby('job_id').agg({'type': lambda x: list(x)}).reset_index()
```

```
# Grouping benefits df by job_id and converting inferred into a list
```

```
inferred_grouped_df = benefits_df.groupby('job_id').agg({'inferred': lambda x: list(x)}).reset_index()
```

```
# Combining the above-created dataframes
```

```
benefits_group_combined_df = benefits_grouped_df.merge(inferred_grouped_df, on='job_id',
how='left')
```

```
# Renaming the column 'type' to 'benefits' for better understanding
```

```
benefits_group_combined_df.rename(columns = {'type':'benefits'}, inplace = True)
```

```
# Grouping industries df by job_id and converting industry_id column into a list
```

```
job_industries_grouped_df = job_industries_df.groupby('job_id').agg({'industry_id': lambda x:
list(x)}).reset_index()
```

```
# Grouping skills df by job_id and converting skill_abr column into a list
```

```
job_skills_grouped_df= job_skills_df.groupby('job_id').agg({'skill_abr': lambda x: list(x)}).reset_index()
```

```
# Merging all datasets with job_postings being the left table
```

```
merged_job_postings_df =
job_postings_df.merge(benefits_group_combined_df,on='job_id',how='left').merge(job_industries_grouped_df,on='job_id',how='left').merge(job_skills_grouped_df,on='job_id',how='left')
```

- Employee counts –
  - This dataset will remain as is since it contains date recorded
  - The only change is that we're creating 2 additional date & time columns, extracted from time\_recorded column to make it easier to query

Snapshot below:

```
# Creating date column to make it easier to query in mongodb
```

```
employee_counts_df['time_recorded_ts'] = pd.to_datetime(employee_counts_df.time_recorded * 1e9)
```

```
employee_counts_df['date_recorded'] = employee_counts_df['time_recorded_ts'].dt.strftime('%Y-%m-%d')
```

These dataframes once created can be exported into MongoDB collections as per the conceptual database design.